

# Homework 6

Submitted By: Javed Roshan

## Part 1: GStreamer

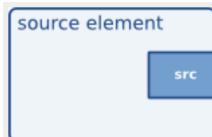
1. In the lab, you used the Nvidia sink `nv3dsink`; Nvidia provides another sink, `nveglglessink`. Convert the following sink to use `nveglglessink`.

```
gst-launch-1.0 v4l2src device=/dev/video0 ! xvimagesink ! nveglglessink
```

```
gst-launch-1.0 v4l2src device=/dev/video0 ! video/x-raw, framerate=30/1 ! videoconvert ! video/x-raw, format=BGR ! nveglglessink
```

2. What is the difference between a property and a capability? How are they each expressed in a pipeline?

An Element is a basic building block for a media pipeline. A ‘source element’ as shown below generates data for use by a pipeline.



A *Pipeline* can be setup by linking a source element with zero or more *Filter*-like elements and finally a *Sink* element. The data will flow through the elements, and this is the basic concept of media handling in *GStreamer*.



*Pads* are the *Element*’s interface to the outside world. Data streams from one element’s source pad to another element’s sink pad. **Capability** is a mechanism that describes the data that can flow through the pad. They are used for compatibility detection, auto-plugging, filtering etc. **Properties** are used to describe extra information of *capabilities*. A property consists of a key (a string) and a value.

Following is an example dump of the capabilities of the “`vorbisdec`” element.

```

Pad Templates:
  SRC template: 'src'
    Availability: Always
    Capabilities:
      audio/x-raw
        format: F32LE
        rate: [ 1, 2147483647 ]
        channels: [ 1, 256 ]

  SINK template: 'sink'
    Availability: Always
    Capabilities:
      audio/x-vorbis

```

As seen in the dump, there are two pads: a source and a sink pad. Both of these pads are *always* available, and both have capabilities attached to them. The sink pad will accept vorbis-encoded audio data, with the media type “audio/x-vorbis”. The source pad will be used to send raw (decoded) audio samples to the next element, with a raw audio media type (in this case, “audio/x-raw”). The source pad will also contain properties for the audio sample rate (rate: 1, 2147483647) and the number of channels (1, 256).

### **3. Explain the following pipeline, that is explain each piece of the pipeline, describing if it is an element (if so, what type), property, or capability. What does this pipeline do?**

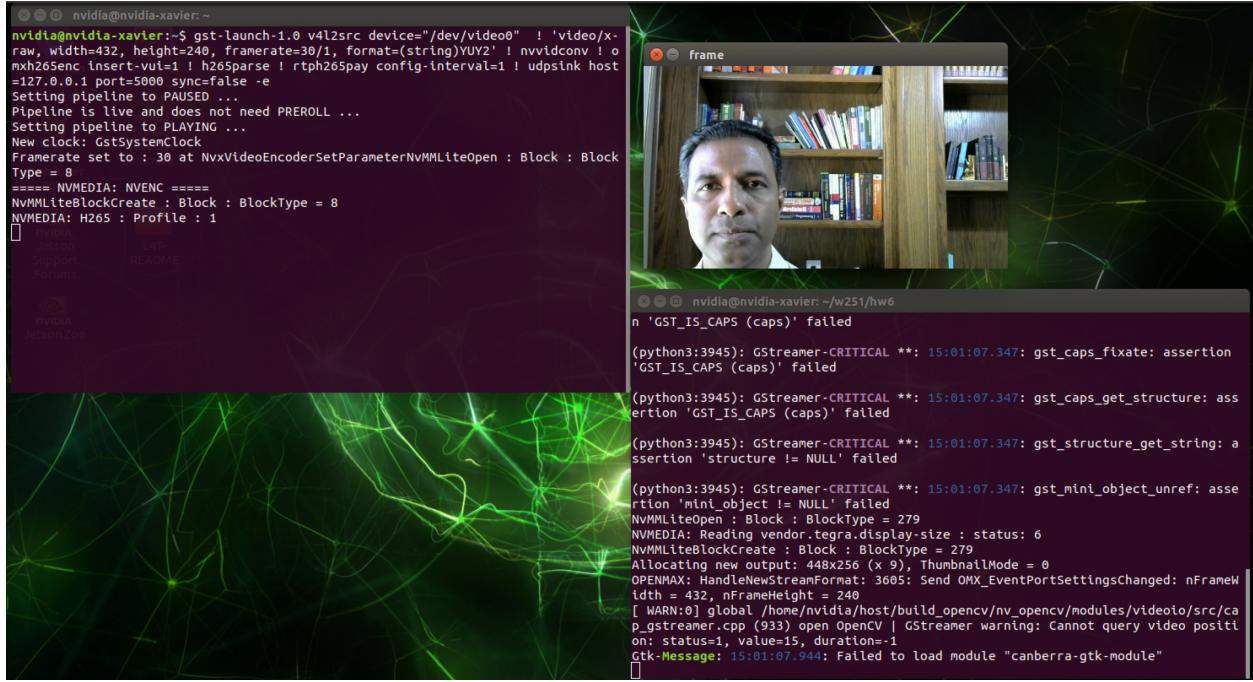
```
gst-launch-1.0 v4l2src device=/dev/video0 ! video/x-raw, framerate=30/1 !
videoconvert ! agingtv scratch-lines=10 ! videoconvert ! xvimagesink sync=false
```

The above pipeline can be de-constructed as follows:

*v4l2src device=/dev/video0* shows that the source video data coming from device 0  
*video/x-raw* defines the capability of the source that has the *framerate* of 30/1  
*videoconvert* is a filter that convert video frames between a great variety of video formats  
*agingtv* ages a video stream in real time, changes the colors and adds scratches and dust, the property *scratch-lines=10* defines the number of scratch lines for this video transformation  
*xvimagesink* renders video frames to a drawable (XWindow) on a local display using the XVideo extension. *sync* is its property when enabled runs the XWindow in synchronous mode



#### 4. GStreamer pipelines may also be used from Python and OpenCV. For example:



#### Publisher:

```
gst-launch-1.0 v4l2src device="/dev/video0" ! 'video/x-raw, width=432, height=240, framerate=30/1, format=(string)YUY2' ! nvvidconv ! omxh265enc insert-vui=1 ! h265parse ! rtph265pay config-interval=1 ! udpsink host=127.0.0.1 port=5000 sync=false
```

#### Listener:

```
Camera Setting: udpsrc port=5000 ! application/x-rtp, media=video, width=640, height=480, payload=96, clock-rate=90000, encoding-name=H265 ! rtph265depay ! h265parse ! omxh265dec ! videoconvert ! appsink
```

```
1 # cam.py: written by Javed Roshan
2 import numpy as np
3 import cv2
4
5 # the index depends on your camera setup and which one is your USB camera.
6 # you may need to change to 1 depending on your local config
7
8 cap=cv2.VideoCapture(0)
9
10 #cap = cv2.VideoCapture(0)
11
12 while(True):
13     # Capture frame-by-frame
14     ret, img = cap.read()
15
16     # Display the resulting frame
17     cv2.imshow('frame', img)
18
19     if cv2.waitKey(1) & 0xFF == ord('q'):
20         break
21
22 # When everything done, release the capture
23 cap.release()
24 cv2.destroyAllWindows()
```

## Part 2: Model optimization & quantization

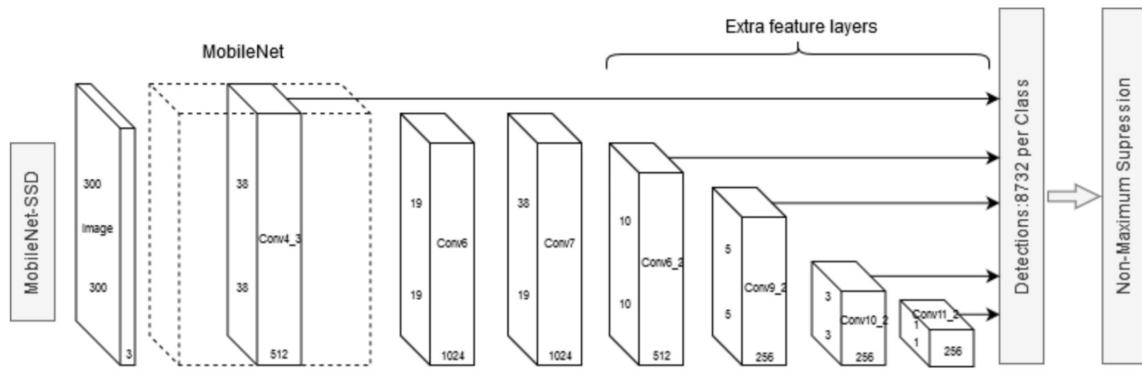
The base model was trained on a total of 2500 images downloaded via

```
python3 open_images_downloader.py --max-images=2500 --class-names  
"Apple,Orange,Banana,Strawberry,Grape,Pear,Pineapple,Watermelon" --data=data3/fruit
```

The base model was downloaded from

```
wget https://nvidia.box.com/shared/static/djf5w54rjvpqocsitzzaandq1m3avr7c.pth -O  
models/mobilenet-v1-ssd-mp-0_675.pth
```

Following is the model architecture:



It was run for 30 epochs with batch size as 4

```
python3 train_ssd.py --data=data3/fruit --model-dir=models/fruit --batch-size=4 --  
epochs=30
```

The data format with some examples is as follows:

ImageID	0026df6bcf93ac5	0026df6bcf93ac5	003be557ebbb9e8c
Source	xclick	Xclick	activemil
LabelName	/m/0kpqd	/m/0kpqd	/m/0kpqd
Confidence	1	1	1
XMin	0.333125	0.488125	0.24875
XMax	0.57625	0.9725	0.326875
YMin	0.270833	0.360833	0.846516
YMax	0.474167	0.999167	0.935028
IsOccluded	1	0	-1
IsTruncated	0	1	-1
IsGroupOf	0	0	-1
IsDepiction	0	0	-1
IsInside	0	0	-1
id	/m/0kpqd	/m/0kpqd	/m/0kpqd
ClassName	Watermelon	Watermelon	Watermelon

The images are linked with the ImageIds and are split into train, validation, & test directories. “fruit” is the category of the images selected for this training.

I was able to convert the native pytorch model into ONNX based model. However, I was unable to convert it into a TF-RT based model. The code and models are available at  
<https://github.com/jroshanucb/py-onnx-tfrt.git>

The training performance is available in the repo in epochs.txt.